

An Equivalent Axiomatization of ZF via a Single *Closure* Schema

Mathematics and a Machine-Verified Rocq/Coq Proof

Vladimir Reshetnikov

Formalization carried out with Claude Code (Rocq 9.0.1)

2026-06-30

Abstract

We study an alternative axiomatization T of Zermelo–Fraenkel set theory obtained by deleting the four “generative” axioms — Pairing, Union, Infinity, and Replacement — and adding in their place a single new schema, *Closure*: the transitive closure of any set under any *set-like* definable relation is a set. Keeping Extensionality, Regularity, Separation, Powerset, and Choice, we show that T is *exactly* $\mathsf{ZF}(\mathsf{C})$. The forward half (the four deleted axioms are theorems of T) is the interesting direction; it rests on one small fact — Powerset gives every set a host — which turns Closure into a fully general principle of collection. The reverse half (every ZF model satisfies Closure) is the textbook transitive-closure recursion. We then describe in tutorial detail a complete machine-checked formalization in the Rocq/Coq proof assistant, organized in four developments: a shallow embedding proving the semantic equivalence in both directions with a “free dependency audit” that certifies *which* axioms each derivation actually uses; a deep embedding of first-order logic that re-proves the forward trade from genuinely syntactic schemas; a sound natural-deduction calculus with the cross-theory corollary $\mathsf{ZF} \vdash \varphi \Rightarrow \mathsf{T} \models \varphi$; and a from-scratch proof of Gödel completeness, compactness for sentence theories, and the forward *syntactic* direction $\mathsf{ZF} \vdash \varphi \Rightarrow \mathsf{T} \vdash \varphi$. The whole development contains no admitted goals and rests only on the standard axioms of classical mathematics. We close by characterizing precisely the one direction that is true but not formalized — the converse $\mathsf{T} \vdash \varphi \Rightarrow \mathsf{ZF} \vdash \varphi$ — and explain why it amounts to formalizing the recursion theorem as a first-order object derivation.

Contents

1	Introduction	2
2	The two theories	3
2.1	The common base	3
2.2	The theory ZF	4
2.3	The theory T	4
3	The forward direction: $\mathsf{T} \vdash \mathsf{ZF}$	5
3.1	The linchpin: Powerset hosts every set	5
3.2	Four derivations, one schema-instance family	5
3.3	Powerset is load-bearing	6
4	The reverse direction: $\mathsf{ZF} \vdash \mathsf{Closure}$	6

5	The equivalence theorem	7
6	Mechanization I: the shallow embedding and the dependency audit	7
6.1	Why Rocq/Coq, and what a “shallow embedding” is	7
6.2	The linchpin and the four theorems, in Coq	8
6.3	The free dependency audit	9
6.4	The reverse construction in Coq	10
6.5	Where classical logic enters	10
7	Mechanization II: closing the first-order gap	10
7.1	Syntax, environments, and Tarski satisfaction	10
7.2	First-order Separation and Closure	11
8	Mechanization III: a proof calculus, soundness, and a cross-theory bridge	12
8.1	A natural-deduction calculus	12
8.2	The equality rule: a genuine correctness fix	13
8.3	Soundness, and the bridge $\text{ZF} \vdash \varphi \Rightarrow \text{T} \models \varphi$	13
9	Mechanization IV: Gödel completeness, compactness, and the syntactic equivalence	14
9.1	The goal	14
9.2	Proof-theory infrastructure	14
9.3	Model existence: the term model and the truth lemma	15
9.4	Lindenbaum and Henkin: building the maximal theory	15
9.5	Compactness: infinite completeness for sentence theories	15
9.6	ZF and T as sentence theories, and $\text{ZF} \vdash \varphi \Rightarrow \text{T} \vdash \varphi$	16
10	What is proven, and the one remaining direction	17
10.1	The converse, characterized precisely	17
10.2	Trust: axioms used, no admitted goals	18
11	Conclusion	18
A	File inventory and build instructions	19

1 Introduction

The usual axiomatization of Zermelo–Fraenkel set theory carries several *generative* axioms whose common job is to assert that certain sets exist: **Pairing** ($\{a, b\}$ exists), **Union** ($\bigcup s$ exists), **Infinity** (an inductive set exists), and **Replacement** (the image of a set under a definable function-class is a set). Each of these is, in isolation, a recipe for collecting a previously described family of sets into one object.

This article studies what happens when we replace all four with a *single* collection principle. Keep the “structural” axioms — Extensionality, Regularity, Separation, Powerset, Choice — and add the following schema.

Closure. For every *set-like* definable binary relation $\prec$ and every set s , there is a set w that contains s and is closed under taking $\prec$ -predecessors. (Such a w is a superset of the transitive closure of s under $\prec$.)

Here “set-like” is the natural smallness condition that makes the principle consistent: each node’s class of immediate $\prec$ -predecessors must already be bounded by some set. Without it, “ $z \prec x$ ” could be “ $z = z$ ”, and the closure of $\{\emptyset\}$ would be the universe.

Call the resulting theory T . The thesis of this article — proved on paper and then in full inside a proof assistant — is:

T and ZF have exactly the same theorems. With the shared axiom of Choice, T and ZFC coincide.

What makes this more than a bookkeeping exercise is the *forward* direction: deriving Pairing, Union, Infinity, and Replacement from the lone Closure schema. We will see that all four are instances of *one* idea, and that the engine which makes them go is not Closure by itself but Closure standing on **Powerset**. The single fact

$$a \subseteq a \implies a \in \mathcal{P}(a)$$

— “Powerset gives every set a host” — is what lets us certify that the relations we care about are set-like, and hence what upgrades Closure from a statement about transitive closures into a fully general principle of collection.

What this article does. There are two interleaved goals. The first is mathematical: state the two theories precisely, prove the equivalence, exhibit the four derivations as one schema-instance family, and isolate the role of Powerset (and, by a pleasant mirror symmetry, the role of Regularity in the reverse direction). The second is to explain, in tutorial style, the *machine-verified* proof we built in the Rocq/Coq proof assistant. That formalization grew well past the original equivalence: it now includes a deep embedding of first-order logic, a sound and *complete* proof calculus (Gödel completeness proved from scratch), a compactness theorem for sentence theories, and the forward syntactic statement $\mathsf{ZF} \vdash \varphi \Rightarrow \mathsf{T} \vdash \varphi$. We will walk through each layer, quote the actual definitions and theorem statements, and explain the design decisions and the one genuine obstacle that remains.

Roadmap. Sections 2–5 are the mathematics: the two theories, the forward direction (with the Powerset remark), the reverse direction, and the equivalence theorem. Sections 6–9 are the formalization: the shallow embedding and the dependency audit (§6), the deep embedding closing the first-order gap (§7), the proof calculus and soundness (§8), and completeness, compactness, and the syntactic equivalence (§9). Section 10 states precisely what is proven and characterizes the single remaining direction. Section 11 reflects on both punchlines.

2 The two theories

We work in the first-order language with one binary relation symbol $\in$ and equality. All schemas allow parameters (free variables other than the displayed ones). Throughout, $\subseteq$ abbreviates the definable inclusion $a \subseteq b := \forall x (x \in a \rightarrow x \in b)$.

2.1 The common base

Both theories contain:

Extensionality

$$\forall a \forall b (\forall x (x \in a \leftrightarrow x \in b) \rightarrow a = b).$$

Separation (schema)

for each formula $\varphi(x, \bar{p})$, $\forall \bar{p} \forall a \exists s \forall x (x \in s \leftrightarrow (x \in a \wedge \varphi(x, \bar{p})))$.

Powerset

$\forall a \exists p \forall x (x \in p \leftrightarrow x \subseteq a)$.

Regularity

every nonempty set has an $\in$ -minimal element.

Choice

(in the C-variants) every set of nonempty sets has a choice function.

2.2 The theory ZF

ZF adds the four generative axioms:

Pairing

$\forall a \forall b \exists p \forall x (x \in p \leftrightarrow (x = a \vee x = b))$.

Union

$\forall s \exists u \forall x (x \in u \leftrightarrow \exists v (x \in v \wedge v \in s))$.

Infinity

there is a set containing $\emptyset$ and closed under $x \mapsto x \cup \{x\}$.

Replacement (schema)

for each functional formula, the image of a set is a set.

2.3 The theory T

T drops all four and adds the single **Closure** schema. Two definitions make it precise. A binary relation $\prec$ is *set-like* when

$$\forall x \exists y \forall z (z \prec x \rightarrow z \in y), \quad (1)$$

i.e. the class of immediate $\prec$ -predecessors of each node is bounded by a set. A set w is *closed under $\prec$ -predecessors above s* when

$$s \subseteq w \wedge \forall u \forall v (u \prec v \wedge v \in w \rightarrow u \in w). \quad (2)$$

The Closure schema asserts one such w for every set-like definable $\prec$ and every s :

$$\underbrace{(\forall x \exists y \forall z (z \prec x \rightarrow z \in y))}_{\prec \text{ set-like}} \implies \forall s \exists w \underbrace{(s \subseteq w \wedge \forall u \forall v (u \prec v \wedge v \in w \rightarrow u \in w))}_{w \text{ contains the transitive closure of } s}. \quad (3)$$

There is one instance of (3) per definable binary relation $\prec$, parameters allowed.

Remark 2.1 (Why w is a *superset*, not the exact transitive closure). We deliberately ask only for a set w *containing* the transitive closure, not for the transitive closure itself. This is harmless: once w exists, Separation carves out the exact transitive closure as a subset of w . Asking only for a superset keeps each Closure instance about an existential “there is a bounding set,” which is what makes the set-like hypothesis exactly the right amount of input.

Remark 2.2 (Choice is a passenger). Choice is stated identically in both theories and is never used in the trade between the generative axioms and Closure. The equivalence therefore reduces to the purely “generative” content: over the common base {Ext, Sep, Pow, Reg}, the four axioms {Pairing, Union, Infinity, Replacement} are interderivable with the single schema Closure.

3 The forward direction: $\mathbf{T} \vdash \mathbf{ZF}$

This is the interesting half. We assume only {Extensionality, Separation, Powerset, Closure} and derive Pairing, Union, Replacement, and Infinity. (Regularity and Choice are not needed here; they are genuine passengers, as the dependency audit in §6 will certify.)

3.1 The linchpin: Powerset hosts every set

Lemma 3.1 (self-membership in the powerset). *For every set a , $a \in \mathcal{P}(a)$.*

Proof. $a \subseteq a$ holds trivially, and by Powerset $x \in \mathcal{P}(a) \leftrightarrow x \subseteq a$; take $x = a$. $\square$

Lemma 3.1 looks too small to matter, but it is the entire engine. Consider any class relation $\prec$ that is *singleton-valued* above each node — i.e. for each x there is at most one z with $z \prec x$, given by some class function $z = G(x)$. Is such a $\prec$ set-like? The set-like condition (1) asks for a set y bounding $\{z : z \prec x\} = \{G(x)\}$. By Lemma 3.1, $G(x) \in \mathcal{P}(G(x))$, so $y = \mathcal{P}(G(x))$ works *for free*. More generally, whenever the predecessor-class above each node is already known to sit inside a definable set, Powerset hands us the bound.

*This is the conceptual heart of the whole equivalence: **Powerset turns “has a definable predecessor” into “is set-like.”** Closure then collects the predecessors. The four generative axioms are four ways of choosing the relation $\prec$ and the seed s .*

3.2 Four derivations, one schema-instance family

Axiom	Relation $\prec$ (read $z \prec x$)	Seed s
Union	$z \in x$ (membership)	the given set s
Replacement	$x \in a \wedge z = F(x)$ (function graph)	the domain a
Pairing	$(x=\emptyset \wedge z=a) \vee (x=\{\emptyset\} \wedge z=b)$	the 2-node seed $\{\emptyset, \{\emptyset\}\}$
Infinity	$z = x \cup \{x\}$ (successor)	$\{\emptyset\}$

In each case the recipe is the same: (i) check the relation is set-like (using Lemma 3.1); (ii) apply Closure to get a closed superset w ; (iii) use Separation to carve the exact set we wanted out of w .

Union is closure under $\in$. The relation $z \prec x := z \in x$ is set-like because the predecessors of x are exactly the elements of x , bounded by x itself ($y = x$). Closure applied to s gives a w with $s \subseteq w$ and w closed under $\in$ -predecessors. Hence every element of every member of s lies in w , and Separation extracts $\{x \in w : \exists v (x \in v \wedge v \in s)\} = \bigcup s$.

Replacement is closure along a function graph. For a class function F and domain a , take $z \prec x := x \in a \wedge z = F(x)$. The predecessors of x form the singleton $\{F(x)\}$ (or $\emptyset$ if $x \notin a$), bounded by $\mathcal{P}(F(x))$ via Lemma 3.1 — this is where Powerset is used. Closure applied to a yields a w containing the whole image $\{F(x) : x \in a\}$, and Separation carves out exactly that image.

Pairing is the subtle one, and it is precisely where Powerset’s host-providing role becomes visible as a *design constraint*. We want $\{a, b\}$ for arbitrary a, b . The naive attempt — one seed node whose predecessors are $\{a, b\}$ — is self-defeating: to certify that relation set-like we would need a set bounding $\{a, b\}$, which is Pairing itself. The fix is to *split a and b across two different seed nodes*. Use the two-node seed $\{\emptyset, \{\emptyset\}\}$ (built from $\emptyset$ and $\mathcal{P}$ alone) and the relation

$$z \prec x := (x = \emptyset \wedge z = a) \vee (x = \{\emptyset\} \wedge z = b).$$

Now each node has a *singleton* predecessor-class — $\{a\}$ above $\emptyset$, $\{b\}$ above $\{\emptyset\}$ — so Lemma 3.1 makes it set-like ($\mathcal{P}(a)$ and $\mathcal{P}(b)$ are the hosts), with no circular appeal to Pairing. Closure of the seed then drags both a and b into one set w , and Separation gives $\{x \in w : x = a \vee x = b\} = \{a, b\}$.

Infinity is closure of $\{\emptyset\}$ under the successor relation $z \prec x := z = x \cup \{x\}$. The successor is unique, so the predecessor-class is again a singleton, hostable by Powerset; the relation is set-like. Closure applied to $\{\emptyset\}$ gives a set w containing $\emptyset$ and closed under successors — an inductive set, which is exactly what Infinity asserts. (Once Pairing and Union are available, the successor $x \cup \{x\}$ is itself a set, so the relation is genuinely definable.)

3.3 Powerset is load-bearing

The four derivations above use Powerset in a single, uniform role: as the source of hosts in Lemma 3.1, the fact that makes singleton-valued (and more generally suitably bounded) relations set-like. Union is the lone exception — its relation $z \in x$ is bounded by x directly, with no powerset needed — which is why the audit in §6 will show Union depending on Separation and Closure *only*.

This is worth stating sharply, because it is easy to read Closure as “the” replacement for the four axioms and miss that it is *Closure together with Powerset* that does the work. Strip Powerset away and the set-like hypothesis of Closure stops being free: there is no longer a uniform host for a singleton-valued relation, and the collection power evaporates. We will see the machine certify this load-bearing role automatically, and — by a pleasing symmetry — certify that Powerset is *idle* in the reverse direction, where Regularity becomes the load-bearing axiom instead.

4 The reverse direction: $\mathbf{ZF} \vdash \text{Closure}$

The reverse direction is the standard transitive-closure construction; the only care needed is in collecting an ω -indexed family of sets into one object, which is exactly the job Replacement and Infinity exist to do.

Fix a set-like relation R and a set s . Define the one-step operator

$$g(t) = t \cup \{u : \exists v \in t, u R v\}.$$

The bracketed predecessor-set is genuinely a set: set-likeness gives, for each v , a bounding set $b(v)$ of $\{u : u R v\}$; Replacement collects $\{b(v) : v \in t\}$, Union flattens it, and Separation carves out the actual predecessors. So $g : V \rightarrow V$ is a definable total operation.

Now iterate: $W_0 = s$ and $W_{n+1} = g(W_n)$, indexing on the *metatheory’s* natural numbers. The closure we want is $w = \bigcup_n W_n$. To produce w as a single set we must turn the meta-indexed family $\{W_n\}$ into an *object*-indexed family. Use Infinity to obtain an inductive set, and inside it the von Neumann numerals $\underline{n}$ ($\underline{0} = \emptyset$, $\underline{n+1} = \underline{n} \cup \{\underline{n}\}$). The assignment $\underline{n} \mapsto W_n$ is a definable class function on the numerals, so Replacement makes its image $\{W_n : n\}$ a set, and Union gives $w = \bigcup_n W_n$.

For this to be well defined we need the numerals to be *distinct*, i.e. $\underline{m} = \underline{n} \Rightarrow m = n$, so that “ $\underline{n} \mapsto W_n$ ” is unambiguous. Distinctness follows from $\underline{m} \in \underline{n}$ whenever $m < n$, together with the fact that *no set is a member of itself* — which is precisely where **Regularity** enters (apply Regularity to $\{x\}$ to refute $x \in x$). Finally, w contains $s = W_0$ and is closed under R -predecessors: if $u R v$ and $v \in w$, then $v \in W_n$ for some n , hence $u \in g(W_n) = W_{n+1} \subseteq w$.

Remark 4.1 (The reverse mirror: Regularity load-bearing, Powerset idle). Notice what powered the reverse construction: Replacement, Union, Infinity (to collect the family) and Regularity (to make the numerals injective). Powerset was *not* used. The machine audit in §6 confirms this, yielding the sharper statement

$$\text{ZF} - \text{Powerset} \vdash \text{Closure}.$$

So the two directions exhibit a clean mirror image: **Powerset is load-bearing forward and idle in reverse; Regularity is idle forward and load-bearing in reverse.**

5 The equivalence theorem

Theorem 5.1 (Equivalence). *Over the common base $\{\text{Extensionality}, \text{Separation}, \text{Powerset}, \text{Regularity}\}$ (and the shared Choice), the theory T (adding Closure) and the theory ZF (adding Pairing, Union, Infinity, Replacement) have exactly the same theorems. In particular T and ZFC are the same theory.*

Proof. $T \vdash ZF$: Pairing, Union, Replacement, and Infinity are Theorems by the four derivations of §3, none of which uses any ZF-only axiom. $ZF \vdash T$: Closure is a Theorem by the construction of §4. Since the two theories share their remaining axioms verbatim, they prove the same sentences. $\square$

The rest of the article is about making Theorem 5.1 — and a good deal more — checkable by machine.

6 Mechanization I: the shallow embedding and the dependency audit

The formalization lives in the `src/SetTheory/` directory of the repository, in four Rocq/Coq source files. This section covers the two that prove the semantic equivalence directly: `Forward.v` and `Reverse.v`.

6.1 Why Rocq/Coq, and what a “shallow embedding” is

We use Rocq/Coq (version 9.0.1). The decisive feature is the **Section/Hypothesis** idiom, which fits our problem like a glove. A Coq **Section** lets us fix an abstract structure and assume axioms about it as local **Hypotheses**; every theorem proved inside is, after the section closes, automatically generalized over the structure *and over exactly the hypotheses it used*. That is precisely the shape of a metatheorem of the form “every model of these axioms satisfies that conclusion.”

The structure is modelled *shallowly*: the universe of sets is an abstract Coq type V , membership is an abstract relation $\text{mem} : V \rightarrow V \rightarrow \text{Prop}$, and the axiom schemas are rendered with the metatheory’s own predicates. For instance Separation quantifies over a Coq predicate $P : V \rightarrow \text{Prop}$:

```
Variable V : Type.
Variable mem : V -> V -> Prop.
```

```

Infix " " := mem (at level 70, no associativity).
Definition Sub (a b : V) : Prop := forall x, x a -> x b.

Hypothesis Extensionality :
  forall a b, (forall x, x a <-> x b) -> a = b.
Hypothesis Separation :
  forall (a : V) (P : V -> Prop),
    exists s, forall x, x s <-> (x a /\ P x).
Hypothesis Powerset :
  forall a, exists p, forall x, x p <-> Sub x a.

```

The Closure schema becomes a quantification over a Coq relation $R : V \rightarrow V \rightarrow \text{Prop}$, with set-likeness defined exactly as in (1):

```

Definition SetLike (R : V -> V -> Prop) : Prop :=
  forall x, exists y, forall z, R z x -> z y.

Hypothesis Closure :
  forall R : V -> V -> Prop, SetLike R ->
    forall s, exists w, Sub s w /\ (forall u v, R u v -> v w -> u w).

```

What a shallow embedding proves, and its faithfulness. Because the schemas range over *all* Coq predicates/relations, the artifact establishes the *second-order* statement: every structure $(V, \in)$ satisfying these axioms also satisfies the derived ones. This is the standard, accepted way to mechanize an equivalence of axiomatizations, and it is strictly stronger than the first-order schematic reading on the model side. It stays faithful to the genuine first-order claim because *every derivation instantiates a schema at one concrete, definable relation* — the membership relation, a function graph, the two-branch pairing relation, the successor relation — exactly the instances a first-order proof would write down. (Section 7 removes even this caveat for the forward direction, by re-proving it with schemas that range over syntactic formulas.)

6.2 The linchpin and the four theorems, in Coq

Lemma 3.1 is a two-line Coq proof, and it is exactly the pivot the informal argument advertised:

```

Lemma Sub_refl : forall a, Sub a a.
Proof. intros a x H. exact H. Qed.

(* THE LINCHPIN: Powerset gives every set a host. *)
Lemma self_in_power : forall a, a power a.
Proof. intro a. apply power_intro. apply Sub_refl. Qed.

```

The four relations of the schema-instance family are declared verbatim:

```

Definition pairRel (a b : V) : V -> V -> Prop :=
  fun z x => (x = emptyset /\ z = a) \/ (x = single_empty /\ z = b).
Definition memRel : V -> V -> Prop := fun z x => z x.
Definition graphRel (F : V -> V) (a : V) : V -> V -> Prop :=
  fun z x => x a /\ z = F x.
Definition succRel : V -> V -> Prop :=
  fun z x => forall t, t z <-> (t x \/ t = x).

```

Each of Pairing, Union, Replacement, Infinity is then proved by the three-step recipe (set-like check $\rightarrow$ Closure $\rightarrow$ Separation). Union is the shortest and shows the pattern cleanly:

```
Theorem Union :
  forall s, exists u, forall x, x u <-> exists v, x v /\ v s.
Proof.
  intro s.
  assert (HSL : SetLike memRel).
  { intro x. exists x. intros z Hz. unfold memRel in Hz. exact Hz. }
  destruct (Closure memRel HSL s) as [w [Hsub Hclosed]].
  exists (sep w (fun x => exists v, x v /\ v s)).
  intro x. split.
  - intro H. exact (sep_elim2 _ _ _ H).
  - intro H. apply sep_intro.
    + destruct H as [v [Hxv Hvs]]. apply (Hclosed x v).
      * unfold memRel. exact Hxv.
      * apply Hsub. exact Hvs.
    + exact H.
Qed.
```

Read it as English: “`memRel` is set-like (host = x itself); apply `Closure` to get a closed superset w of s ; separate out the genuine union elements; membership in either direction is immediate from closure/Separation.”

The Pairing proof is where the two-node seed appears and where the set-like check does real work: it splits on whether the node is $\emptyset$ or $\{\emptyset\}$ and supplies $\mathcal{P}(a)$ or $\mathcal{P}(b)$ as the host accordingly — the Coq transcription of the “split a and b across different nodes” argument from §3.

6.3 The free dependency audit

Here is the part that turns the “Powerset is load-bearing” remark from rhetoric into a machine certificate. When a Coq `Section` closes, each theorem is generalized over precisely the hypotheses it referenced; the trailing `Check` commands print those generalized types, so we can simply *read off* which axioms each derivation used:

```
End ClosureEquivalence_Forward.
Check Pairing.
Check Union.
Check Replacement.
Check Infinity.
```

The printed signatures certify:

- **Union** depends on *Separation and Closure only* — not Powerset, not Extensionality, not the nonempty-domain assumption. (Its relation $z \in x$ is self-bounding.)
- **Replacement** depends on *Separation, Powerset, and Closure* — Powerset in exactly its host-providing role.
- **Pairing** and **Infinity** additionally need *Extensionality* and a *nonempty domain* (to build $\emptyset$ and to tell $\emptyset \neq \{\emptyset\}$ apart).
- **Regularity** is referenced by *none* of the four — a genuine passenger, just as §3 claimed.

The same audit on `Reverse.v` (`Check Closure_holds`) shows `Closure_holds` depending on a nonempty domain, Extensionality, Separation, Pairing, Union, Infinity, Replacement, and **Regularity** — but **not Powerset**. This is the machine confirming the sharper $\text{ZF} - \text{Powerset} \vdash \text{Closure}$ and the forward/reverse mirror of Remark 4.1: Powerset load-bearing forward and idle in reverse, Regularity idle forward and load-bearing in reverse.

6.4 The reverse construction in Coq

`Reverse.v` transcribes §4 directly. The meta-level iteration is an ordinary Coq `Fixpoint` on `nat`:

```
Fixpoint iterate (g : V -> V) (s : V) (n : nat) : V :=
  match n with 0 => s | S k => g (iterate g s k) end.
```

The non-self-membership fact — the load-bearing use of Regularity — and the injectivity of the numerals it powers are:

```
Lemma no_self_mem : forall x, ~ x x. (* from Regularity on {x} *)
Lemma onat_inj : forall i j, onat i = onat j -> i = j.
```

`onat_inj` is what makes the map “numeral $\mapsto$ iterate” well defined, so Replacement can collect $\{W_n\}$ into a set and Union flatten it. The headline theorem is then

```
Theorem Closure_holds :
  forall R : V -> V -> Prop, SetLike R ->
    forall s, exists w, Sub s w /\ (forall u v, R u v -> v w -> u w).
```

with $w = \bigcup_n W_n$ realized as `ounion (imageR (Ffun R HSL s) Inf)` — the object union of the Replacement-image, over the inductive set `Inf`, of the numeral-to-iterate map `Ffun`.

6.5 Where classical logic enters

Both files import `ClassicalEpsilon`. It is used to *package* the existential axioms as operations — turning “ $\exists p, \dots$ ” (Powerset) into a function `power : V -> V` with a specification lemma, via `constructive_indefinite_description` — and, in `Reverse.v`, to extract a bounding function from set-likeness and to index the iteration. Set theory is classical anyway and Choice rides along, so this is exactly the intended metatheoretic setting.

7 Mechanization II: closing the first-order gap

The shallow embedding’s one honest caveat is that its schemas range over arbitrary Coq predicates rather than over first-order *formulas*. `Deep.v` removes that caveat for the forward direction by building a genuine deep embedding of the first-order language of set theory and re-proving the trade from schemas that quantify over syntactic formulas.

7.1 Syntax, environments, and Tarski satisfaction

A formula is an element of an inductive type, with De Bruijn indices for variables (so a variable is a `nat` counting binders outward, and α -equivalence is syntactic equality):

```
Inductive form : Type :=
```

```

| fMem : nat -> nat -> form
| fEq  : nat -> nat -> form
| fBot : form
| fImp : form -> form -> form
| fAnd : form -> form -> form
| fOr  : form -> form -> form
| fAll : form -> form
| fEx  : form -> form.

```

An *environment* $e : \text{nat} \rightarrow V$ assigns a domain element to each variable. `scons d e` pushes a new binding `d` at index 0 and shifts the rest — this is how a quantifier extends the environment under its binder. Tarski satisfaction is then a straightforward structural recursion:

```

Definition scons (d : V) (e : nat -> V) : nat -> V :=
  fun n => match n with 0 => d | S k => e k end.

Fixpoint Sat (e : nat -> V) (f : form) {struct f} : Prop :=
  match f with
  | fMem i j => (e i) = (e j)
  | fEq i j  => e i = e j
  | fBot     => False
  | fImp a b => Sat e a -> Sat e b
  | fAnd a b => Sat e a /\ Sat e b
  | fOr a b  => Sat e a \/ Sat e b
  | fAll a   => forall d, Sat (scons d e) a
  | fEx a    => exists d, Sat (scons d e) a
  end.

```

Two infrastructural lemmas carry the whole development. `Sat_ext` says satisfaction depends only on the environment’s values on the variables that occur; `Sat_rename` relates satisfaction of a renamed formula to satisfaction under the renamed environment:

$$\text{Sat } e (\text{rename } r f) \leftrightarrow \text{Sat } (e \circ r) f.$$

Renaming is the deep-embedding analogue of variable substitution; because our signature is *purely relational*, there are no function symbols and hence no compound terms — a “term” is just a variable — so substitution of a variable for a variable is exactly a renaming. This single observation keeps the entire metatheory free of a general substitution operation, which is usually the most painful part of a deep embedding.

7.2 First-order Separation and Closure

Now the schemas range over formulas. A formula `phi` with free variable 0 denotes a unary predicate (its other free variables are parameters, read from the environment `e`); a formula `psi` with free variables 0, 1 denotes a binary relation `relOf psi e`:

```

Hypothesis SeparationF0 :
  forall (phi : form) (e : nat -> V) (a : V),
    exists s, forall x, x s <-> (x a /\ Sat (scons x e) phi).

Definition relOf (psi : form) (e : nat -> V) : V -> V -> Prop :=
  fun z x => Sat (scons z (scons x e)) psi.

```

```

Hypothesis ClosureF0 :
  forall (psi : form) (e : nat -> V),
    SetLike (relOf psi e) ->
      forall s, exists w, Sub s w /\ (forall u v, relOf psi e u v -> v w -> u w).

```

Deep.v then re-derives Pairing, Union, first-order Replacement, and Infinity from *these* schemas. The new obligation, compared to the shallow version, is that each relation used must be exhibited as a concrete `form` and its `Sat`-meaning verified. For the simple cases the defining formula is short; for Replacement, expressing “ $\exists x \in a, \psi(y, x)$ ” as a formula requires shuffling De Bruijn indices, which is exactly what `Sat_rename` is for. The upshot is a *formal* certificate — not a meta-remark — that the schema-trade uses only first-order definable instances:

```

Theorem Pairing      : forall a b, exists p, forall x, x p <-> (x = a \/ x = b).
Theorem Union       : forall s, exists u, forall x, x u <-> exists v, x v /\ v s.
Theorem ReplacementF0 : (* image of a set under a formula-defined functional relation *)
Theorem Infinity    : (* an inductive set *)

```

Why only the forward direction deepens. The forward trade ports to the deep setting cheaply because every relation it uses is a short first-order formula. The reverse direction does *not*: its essential use of Replacement collects the family $\{W_n\}$ via the map $\underline{n} \mapsto W_n$, and that map is first-order definable only through the *recursion theorem*. `Reverse.v` deliberately sidesteps the recursion theorem by iterating on the meta-level `nat`; rendering the same construction first-order would mean formalizing the recursion theorem’s defining formula — a heavy object-level construction. So the genuine first-order content of the reverse direction *is* the recursion theorem, and we leave it at the (already verified) shallow level. We return to this point in §10.

8 Mechanization III: a proof calculus, soundness, and a cross-theory bridge

So far everything has been semantic (about models). `Deep.v` also builds a *syntactic* layer: an actual proof calculus over `form`, a soundness theorem, and a first bridge between the two theories at the level of provability.

8.1 A natural-deduction calculus

`Prov G a` means “*a* is derivable from the context *G* (a list of formulas).” The calculus is classical natural deduction:

```

Inductive Prov : list form -> form -> Prop :=
| P_ass      : forall G a, In a G -> Prov G a
| P_impI     : forall G a b, Prov (a :: G) b -> Prov G (fImp a b)
| P_impE     : forall G a b, Prov G (fImp a b) -> Prov G a -> Prov G b
| P_botE     : forall G a, Prov G fBot -> Prov G a
| P_lem      : forall G a, Prov G (fOr a (fImp a fBot))
| P_andI     : forall G a b, Prov G a -> Prov G b -> Prov G (fAnd a b)
| P_andE1    : forall G a b, Prov G (fAnd a b) -> Prov G a
| P_andE2    : forall G a b, Prov G (fAnd a b) -> Prov G b

```

```

| P_orI1   : forall G a b, Prov G a -> Prov G (fOr a b)
| P_orI2   : forall G a b, Prov G b -> Prov G (fOr a b)
| P_orE    : forall G a b c, Prov G (fOr a b) -> Prov (a :: G) c -> Prov (b :: G) c
            -> Prov G c
| P_allI   : forall G a, Prov (map (rename S) G) a -> Prov G (fAll a)
| P_allE   : forall G a k, Prov G (fAll a) -> Prov G (rename (inst k) a)
| P_exI    : forall G a k, Prov G (rename (inst k) a) -> Prov G (fEx a)
| P_exE    : forall G a c, Prov G (fEx a) -> Prov (a :: map (rename S) G) (rename S
            c) -> Prov G c
| P_eqRef1 : forall G k, Prov G (fEq k k)
| P_eqElim : forall G i j a,
            Prov G (fEq i j) -> Prov G (rename (inst i) a) -> Prov G (rename (inst j) a).

```

The De Bruijn discipline shows up in the quantifier rules. In `P_allI` (universal introduction), to prove $\forall a$ we shift the context outward by `rename S` — index 0 becomes the fresh “eigenvariable” bound by the new $\forall$, and the old context variables move up by one. In `P_allE` (universal elimination), instantiating $\forall a$ at variable k is the renaming `rename (inst k) a`, where `inst k` maps the bound index 0 to k and decrements the rest. The existential rules are dual. Crucially, instantiation is *just renaming* — the purely relational signature again paying off — so the same `rename/Sat_rename` machinery handles quantifiers with no separate substitution.

8.2 The equality rule: a genuine correctness fix

The two equality rules deserve a word, because getting them right was a real correction during development. The reflexivity rule `P_eqRef1` is obvious. For the elimination rule, a tempting but *too weak* design is a “replace-the-variable-everywhere” congruence: from $i = j$ and φ , infer φ with every occurrence of i rewritten to j . That rule cannot even derive *symmetry* of equality: to get $j = i$ from $i = j$ you would rewrite i to j in the formula $j = i$, but $j = i$ contains no i in the right slot to key on, and the rewrite erases the very occurrence you need.

The fix is the proper **Leibniz** rule `P_eqElim`: a formula a is treated as a property with a hole at De Bruijn index 0; from $i = j$ and “ a with the hole filled by i ” (`rename (inst i) a`) we infer “ a with the hole filled by j ” (`rename (inst j) a`). This single rule yields symmetry, transitivity, and full congruence, and — as we verify — it is sound. Without it the completeness theorem of §9 would have been unreachable, since the term model is a quotient by provable equality and needs all the equality reasoning. A small example of the rule earning its keep: it is exactly what makes the term-model congruence lemmas go through.

8.3 Soundness, and the bridge $\mathbf{ZF} \vdash \varphi \Rightarrow \mathbf{T} \models \varphi$

Soundness says the calculus only proves things that are true in every model satisfying the context:

```

Theorem soundness :
  forall G a, Prov G a -> forall e, (forall x, In x G -> Sat e x) -> Sat e a.

```

The proof is one case per rule; the quantifier and equality cases are where `Sat_rename` and `Sat_ext` are spent.

Now encode the ZF axioms as closed formulas (`Ext_form`, `Pair_form`, ..., `Reg_form`, with Separation and Replacement as schemas over `form`), collect them into a predicate `ZF_ax`, and define provability from the ZF axioms:

```

Definition ZFprov (phi : form) : Prop :=
  exists G, (forall x, In x G -> ZFax x) /\ Prov G phi.

```

```

Theorem ZF_provable_holds_in_T :
  forall phi, ZFprov phi -> forall e, Sat e phi.

```

Because the ambient section assumes the T-axioms, this theorem says: *in any model of the Closure axiomatization T, every ZF-provable formula holds*, i.e. $\text{ZF} \vdash \varphi \Rightarrow \text{T} \models \varphi$. It is the marriage of syntactic soundness with the forward semantic equivalence: each encoded ZF axiom is checked to hold in the T-model (using the derived **Pairing**, **Union**, **ReplacementF0**, **Infinity** for the four traded ones, and the T-hypotheses for the shared ones), and then **soundness** propagates truth through any ZF derivation.

The symmetric statement $\text{T} \vdash \varphi \Rightarrow \text{ZF} \models \varphi$ needs “every ZF-model satisfies **ClosureF0**” — the deep reverse direction, blocked at the recursion theorem as discussed.

9 Mechanization IV: Gödel completeness, compactness, and the syntactic equivalence

The fourth file, **Completeness.v**, proves the converse of soundness — **Gödel completeness** — from scratch, then leverages it to obtain compactness for sentence theories and the forward *syntactic* equivalence $\text{ZF} \vdash \varphi \Rightarrow \text{T} \vdash \varphi$. (It builds on **Deep.v** through Coq’s namespacing: `coqc -Q . SetTheory Deep.v` then `coqc -Q . SetTheory Completeness.v`.)

9.1 The goal

```

Theorem completeness :
  forall G phi,
    (forall (Dom : Type) (m : Dom -> Dom -> Prop) (v : nat -> Dom),
      (forall g, In g G -> Sat Dom m v g) -> Sat Dom m v phi) ->
    Prov G phi.

Corollary prov_iff_valid :
  forall G phi,
    Prov G phi <->
    (forall (Dom : Type) (m : Dom -> Dom -> Prop) (v : nat -> Dom),
      (forall g, In g G -> Sat Dom m v g) -> Sat Dom m v phi).

```

In words: *a formula is provable in the calculus if and only if it is valid in every model*. The “only if” is soundness; the “if” is the content of completeness. The standard route is Henkin’s: show that a *consistent* set of formulas has a model (then contrapose). We outline the construction as it appears in the file.

9.2 Proof-theory infrastructure

The file first develops the usual admissible rules: weakening, the deduction theorem, proof by contradiction and double-negation elimination, a notion of consistency $\text{Con } G := \sim \text{Prov } G \text{ fBot}$, the Lindenbaum extension step **Con_cons_or** (a consistent set stays consistent when extended by a formula or its negation), the equality kit (symmetry/transitivity/congruence — available only

because the equality rule was corrected to the Leibniz `P_eqElim`), the admissibility of renaming `Prov_rename`, and cut `Prov_cut`.

9.3 Model existence: the term model and the truth lemma

The heart is `model_exists`: an abstract *maximal-consistent Henkin* theory `T` (one that decides every formula and provides a witness for every existential) has a model. The model is the *quotient term model*. Since terms are just variables (`nat`), the domain ought to be variables modulo provable equality `ceq i j`. To get an honest Coq type rather than a setoid, we pick *canonical representatives* with Hilbert's ε :

```
Definition rep (i : nat) : nat := epsilon (inhabits 0) (fun j => ceq i j).
Definition D : Type := { n : nat | rep n = n }.          (* canonical reps *)
Definition memD (x y : D) : Prop := cmem (proj1_sig x) (proj1_sig y).
```

The linchpin is the **truth lemma**: in this model, satisfaction of a formula coincides with its membership in the theory `T`. It is proved by strong induction on formula *size* (`fsize`), because the quantifier cases recurse on an instantiated formula — and instantiation is a renaming, which preserves size:

```
Lemma truth :
  forall n a, fsize a <= n ->
    forall s : nat -> nat,
      Sat D memD (fun i => mkD (s i)) a <-> T (rename s a).
```

The atomic cases use the congruence of `cmem/ceq` under the representative map; the connective cases use that `T` is maximal-consistent (it respects $\rightarrow$, $\wedge$, $\vee$, $\neg$); the quantifier cases use that `T` is Henkin (an existential is in `T` iff some witness instance is, and dually for $\forall$).

9.4 Lindenbaum and Henkin: building the maximal theory

To apply `model_exists` we must *construct* a maximal-consistent Henkin extension of any consistent set. This needs an enumeration of all formulas; we build one from Cantor's pairing function and prove it surjective:

```
Definition Enum (n : nat) : form := decode (S n) n.
Lemma Enum_surj : forall f, exists n, Enum n = f.
```

Then a chain `chain G0 n` processes formulas one at a time: at each step it decides the next formula (adding it or its negation, by `Con_cons_or`) and, when it adds an existential, also adds a *fresh-witness* instance (the eigenvariable and Henkin-witness lemmas, from `Prov_rename` plus a freshness analysis of free variables). The limit theory `TL` is shown to satisfy all five hypotheses of `model_exists` — consistent, complete, deductively closed, and Henkin for both $\exists$ and $\forall$. Combining: `model_of_con` (every consistent set has a model), and contraposing gives **completeness**, hence `prov_iff_valid`.

9.5 Compactness: infinite completeness for sentence theories

Completeness as stated is about a *finite* context `G` (a list). To reach genuinely infinite theories — such as ZF and `T`, whose Separation/Closure schemas have infinitely many instances — we need the

compactness-style lift: if every model of an infinite theory B satisfies φ , then B proves φ from a finite subset.

The obstacle is Henkin freshness. The witness-adding step needs a variable *fresh for the whole theory*; for an infinite theory that mentions all variables, there is none. The resolution is to restrict to *sentence* theories — theories all of whose axioms are sentences (no free variables):

```

Definition Sentence (f : form) : Prop := forall n, ~ Free n f.
Definition Sentences (B : form -> Prop) : Prop := forall f, B f -> Sentence f.
Definition BProv (B : form -> Prop) (G : list form) (phi : form) : Prop :=
  exists Gb, (forall x, In x Gb -> B x) /\ Prov (Gb ++ G) phi.

```

A witness fresh for the *finite* list of formulas added so far is then automatically fresh with respect to the sentence base B — sentences contribute no free variables to clash with. This yields:

```

Theorem model_of_BCon :
  forall B L0, Sentences B -> BCon B L0 ->
    exists Dom m v, (forall g, B g -> Sat Dom m v g)
      /\ (forall g, In g L0 -> Sat Dom m v g).

Theorem completeness_inf :
  forall B psi, Sentences B -> Sentence psi ->
    (forall Dom m v, (forall g, B g -> Sat Dom m v g) -> Sat Dom m v psi) ->
    BProv B nil psi.

```

and, as an immediate corollary, the model-theoretic characterization of deductive equivalence:

```

Theorem theory_equiv :
  forall B1 B2, Sentences B1 -> Sentences B2 ->
    (forall Dom m v, (forall g, B1 g -> Sat Dom m v g) <-> (forall g, B2 g -> Sat Dom
      m v g)) ->
    forall psi, Sentence psi -> (BProv B1 nil psi <-> BProv B2 nil psi).

```

`theory_equiv` is the abstract engine behind the whole article: *two sentence theories with the same models prove the same sentences*.

9.6 ZF and T as sentence theories, and $\mathbf{ZF} \vdash \varphi \Rightarrow \mathbf{T} \vdash \varphi$

To feed ZF and T into this machinery we encode them as sentence theories. Each axiom is universally closed by `seal` (close every free variable with $\forall$), so the parametrized schema instances become genuine sentences; crucially, the Closure schema is encoded as a closed formula `Closure_form` whose `Sat`-meaning is bridged back to the abstract `SetLike`/closure statement (`bridge_Closure` and friends):

```

Definition seal (f : form) : form := closeN (bound f) f.      (* universal closure *)
Definition Tax_s (f : form) : Prop := (* sealed Ext, Reg, Pow, Sep-schema,
  Closure-schema *)
Definition ZFax_s (f : form) : Prop := (* sealed Ext, Reg, Pow, Pair, Union, Inf,
  Sep-schema, Repl-schema *)

```

The semantic content of the forward equivalence is repackaged as: *every model of T is a model of ZF*.

```

Lemma Tmodel_sat_ZF :
  forall (V : Type) (mem : V -> V -> Prop) v,
    (forall g, Tax_s g -> Sat V mem v g) -> (forall g, ZFax_s g -> Sat V mem v g).

```

The proof extracts the abstract axioms from a T-model through the `bridge_*` lemmas, then reuses `Deep.v`'s derived satisfaction facts (`sat_Pair`, `sat_Union`, `sat_Inf`, `sat_Repl`) to discharge the four traded ZF axioms. Soundness plus `completeness_inf` then deliver the forward *syntactic* direction:

```

Theorem ZF_implies_T :
  forall phi, Sentence phi -> BProv ZFax_s nil phi -> BProv Tax_s nil phi.

```

Everything ZF proves, T proves. Read against `theory_equiv`, this is one half of the statement that ZF and T are the same theory — and the half that goes through cleanly.

10 What is proven, and the one remaining direction

Statement	Status	File
$T \models \text{Pairing, Union, Replacement, Infinity}$	machine-checked	<code>Forward.v</code>
$ZF \models \text{Closure (set-like relations)}$	machine-checked	<code>Reverse.v</code>
forward trade from <i>first-order</i> schemas	machine-checked	<code>Deep.v</code>
sound ND calculus; $ZF \vdash \varphi \Rightarrow T \models \varphi$	machine-checked	<code>Deep.v</code>
truth lemma / model existence	machine-checked	<code>Completeness.v</code>
completeness $\text{Prov } G \varphi \Leftrightarrow G \models \varphi$	machine-checked	<code>Completeness.v</code>
compactness ($B \models \varphi \Rightarrow B \vdash \varphi$, sentences)	machine-checked	<code>Completeness.v</code>
same-model sentence theories are deductively equal	machine-checked	<code>Completeness.v</code>
$ZF \vdash \varphi \Rightarrow T \vdash \varphi$	machine-checked	<code>Completeness.v</code>
$T \vdash \varphi \Rightarrow ZF \vdash \varphi$	needs the recursion theorem	—

10.1 The converse, characterized precisely

The one direction we have *not* formalized is the converse syntactic statement $T \vdash \varphi \Rightarrow ZF \vdash \varphi$. It is true. By `theory_equiv` it would follow from “every ZF-model satisfies `ClosureF0`,” the dual of `Tmodel_sat_ZF`. That is exactly the *deep reverse* direction — and it is genuinely hard, for a reason that has been visible since §7.

`Reverse.v` proves `Closure` using a *second-order* ingredient: the iteration W_n lives on the metatheory’s `nat`, and the family $\{W_n\}$ is collected by an external recursion. A first-order ZF-model is not handed such a meta-recursion; it must *build* the family internally, as a set-coded function, using only first-order ZF derivations. Doing that is precisely the content of the **recursion theorem** on ω :

- encode functions as sets of Kuratowski pairs;

- formalize the “ n -approximation” of the iteration and prove existence and uniqueness of approximations by ω -induction, *inside* the object calculus **Prov**;
- take the union of approximations to obtain the total iteration function, then its range.

None of this is conceptually mysterious — it is standard textbook set theory — but rendering it as first-order *object derivations* (rather than meta-level Coq recursion) is a self-contained and heavy formalization in its own right. It is the honest boundary of the present development: we prove the converse *semantically* (it is implicit in **Reverse.v**’s shallow model argument) but not *syntactically inside Prov*, because the syntactic content of the reverse direction simply *is* the recursion theorem.

10.2 Trust: axioms used, no admitted goals

Running **Print Assumptions** on the headline theorems reports only the standard axioms of classical mathematics, and nothing else:

```
classic                                (* excluded middle in Prop *)
constructive_indefinite_description    (* Hilbert epsilon / description *)
functional_extensionality
propositional_extensionality
proof_irrelevance
```

There is no **Admitted** and no custom **Axiom** anywhere in the four files. These five are exactly the principles one expects when mechanizing classical set theory with a quotient term model: classical logic, a description operator to turn existence axioms into operations, and the extensionality/irrelevance principles that make the quotient an honest type. In particular, the metatheory is *not* assumed to be stronger than ZF — the development does not, for instance, rely on a universe of ZF-sets supplied by a library.

11 Conclusion

Two punchlines, one mathematical and one about formalization.

The mathematics. Four of ZF’s axioms — Pairing, Union, Infinity, Replacement — are not four independent ideas but four instances of *one*: collect the predecessors of a seed under a set-like relation. The single Closure schema packages that idea, and over the structural base it is interderivable with all four. The non-obvious lesson is *where the power comes from*: not from Closure alone but from Closure resting on Powerset, through the one-line fact that every set is a member of its own powerset. That fact makes “has a definable predecessor” imply “is set-like” for free, which is what turns a statement about transitive closures into general collection. The forward/reverse symmetry is a small gem: Powerset is the load-bearing axiom forward and idle in reverse, while Regularity is idle forward and load-bearing in reverse — and the proof assistant *certifies* this for us, by reading off exactly which hypotheses each derivation discharged.

The formalization. We did not stop at the equivalence. The development ascends from a shallow model argument, to a deep embedding with genuinely first-order schemas, to a sound and *complete* proof calculus built from scratch (truth lemma, Lindenbaum/Henkin, Gödel completeness), to compactness for sentence theories, to the abstract theorem that same-model sentence theories are deductively equivalent — of which “ZF and T prove the same sentences” is an instance, with its

forward half $\text{ZF} \vdash \varphi \Rightarrow \text{T} \vdash \varphi$ now a machine-checked corollary. The single remaining direction is characterized exactly: it is the recursion theorem rendered as a first-order object derivation, the deep reverse obstacle, and nothing less. Everything else is done, with no admitted goals and only the standard classical axioms — a complete and honest account of what it takes to certify that a slogan (“one closure schema”) really is ZF in disguise.

A File inventory and build instructions

The development is four Rocq/Coq files under `src/SetTheory/`.

File	Contents
<code>Forward.v</code>	shallow embedding; <code>{Ext, Sep, Pow, Closure}</code> derive <code>Pairing/Union/Replacement/Infinity</code> ; the linchpin <code>self_in_power</code> ; dependency-audit <code>Checks</code> .
<code>Reverse.v</code>	shallow embedding; ZF derives <code>Closure</code> via the transitive-closure recursion; <code>no_self_mem</code> , <code>onat_inj</code> ; audit shows <code>Powerset</code> unused.
<code>Deep.v</code>	deep embedding (<code>form</code> , <code>Sat</code> , <code>rename</code> , <code>Sat_rename</code>); first-order <code>SeparationF0/ClosureF0</code> ; forward trade re-proved; ND calculus <code>Prov</code> ; <code>soundness</code> ; <code>ZF_provable_holds_in_T</code> .
<code>Completeness.v</code>	from-scratch Gödel completeness/ <code>prov_iff_valid</code> ; <code>completeness_inf</code> (compactness); <code>theory_equiv</code> ; <code>ZF_implies_T</code> .

`Forward.v`, `Reverse.v`, and `Deep.v` are independent (no inter-file `Require`) and need only the standard library. `Completeness.v` `Requires` `Deep` (hence the namespace flag) and additionally uses the standard classical/extensionality modules.

```
# developed against Rocq 9.0.1
coqc Forward.v
coqc Reverse.v
coqc Deep.v
# Completeness.v builds on Deep via the SetTheory namespace:
coqc -Q . SetTheory Deep.v
coqc -Q . SetTheory Completeness.v
```